

Project 1: Leap Frog Pseudo-Random Number Generator

Siddhartha Datta Roy and Vishnu Rajeevan

October 4, 2012

Implementation

We have implemented the leap frog method in kernel code. Executing the leap frog method in kernel space versus userspace leads to timing and performances boosts.

The basis of our implementation depends on each thread writing to the proc file before any reads are done. Each time a write is done, the code checks if the thread's parent process is already stored as a process that is already requesting random numbers. If it is not, the thread's parent process is saved in a linked list. The linked list then points to a linked list of all the threads, that the process owns, that are requesting random numbers. When a random number is requested, the thread is cross referenced against the aforementioned linked list. If the process/thread exists, then, based on the seeding and previous random numbers given to that process/thread, a new random number is outputted. If the thread has not previously requested a random number, then a default value is given. The random number generator is just seeded with the default value of A.

There are a few limitations with this module. First, if the thread is reading before writing, it is always seeded with the same number. Thus, the random numbers generated are always the same. Also, another limitation is that all threads have to write before reading in a process. A thread cannot mix reading and writing when all threads in the same process require the same seed.

One design decision we made is to check if the seed input is a long in the user code rather than the kernel code. This decision was made so that if a non numerical value was given the module would either have to error or choose a default value for the seed. We did not want to have to choose a default value and wanted to error. This is more efficient to do in the user space as there is no context switch that is required when the error is given in the user space.